

Nuo (Tony) Chen, chen2886@purdue.edu

Lab2 Report

- Git: <https://github.com/Chen2886/CS390-Lab2>
- List of resources:
 - o <https://www.cs.toronto.edu/~kriz/cifar.html>
 - o <https://keras.io/api/datasets/cifar100/>
 - o <https://www.tensorflow.org/datasets/catalog/cifar100>
 - o <https://www.tensorflow.org/tutorials/images/cnn>
 - o <https://stackoverflow.com/questions/49079115/valueerror-negative-dimension-size-caused-by-subtracting-2-from-1-for-max-pool>
 - o https://medium.com/@birdortyedi_23820/deep-learning-lab-episode-2-cifar-10-631aea84f11e
 - o
- Part of lab completed
 - o [15] Neural net models.
 - [5] Standard ANN (this can be from lab 1).
 - [10] CNN
 - o [55] Accuracy
 - [10] Get an MNIST digit accuracy of 99% or higher.
 - [10] Get an MNIST fashion accuracy of 92% or higher.
 - [11] Get a cifar 10 accuracy of 70% or higher.
 - [12] Get a cifar 100 course accuracy of 50% or higher.
 - [12] Get a cifar 100 fine accuracy of 35% or higher.
 - [+3] EC: Have a cifar 100 fine accuracy of 40% or higher.
 - [+5] EC: Have a cifar 100 fine accuracy of 45% or higher.
 - o [10] Pipeline & misc.
 - [2] Code the pipeline to be able to use cifar-10
 - [2] Code the pipeline to be able to use cifar-100 fine
 - [2] Code the pipeline to be able to use cifar-100 course
 - [2] Use matplotlib to create a bar graph showing your accuracy with the standard ANN for each dataset. Save this graph as an image and name the file ANN_Accuracy_Plot.pdf.
 - [2] Use matplotlib to create a bar graph showing your accuracy with the standard CNN for each dataset. Save this graph as an image and name the file CNN_Accuracy_Plot.pdf.
 - [+1] EC: Add the option to preprocess the image using random crops for data augmentation. This option will work best on the cifar datasets.
 - [+1] EC: Code in an option to save your network weights to a file for your CNN and ANN, as well as load them. You should be able to run the network again without training.
 - o [20] Report.

- How is a CNN superior to standard ANNs for image processing?
 - o CNN is better because it takes account the special relation.
- Why do we sometimes use pooling in CNNs?
 - o Pooling is to reduce the runtime of CNN, otherwise it will be a long time to train a CNN.
- Why do you think the cifar datasets are harder than mnist?
 - o Cifar images are colored, there are more classes.
- Increases the accuracy of your CNN
 - o Dividing the data by 255 increase the accuracy significantly.
 - o I modeled my CNN after VGG. Dropout layer increased accuracy
 - o Flattened layer and Dense layer at the end increased accuracy
 - o Adding batch normalization increased accuracy
- Hyperparameters:
 - o LR: 0.0008
 - o Activation function: Mostly relu, softmax at the end
 - o Adam optimizer
 - o Categorical cross entropy loss function
 - o 10 Epochs
 - o Drop rate: 0.2
- MNIST F (CNN):
 - o Epoch 1/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.6439 - accuracy: 0.7693
 - o Epoch 2/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.3409 - accuracy: 0.8730
 - o Epoch 3/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.2891 - accuracy: 0.8933
 - o Epoch 4/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.2666 - accuracy: 0.9003
 - o Epoch 5/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.2514 - accuracy: 0.9076
 - o Epoch 6/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.2396 - accuracy: 0.9108
 - o Epoch 7/10
 - o 1875/1875 [=====] - 14s 7ms/step - loss: 0.2288 - accuracy: 0.9144
 - o Epoch 8/10

- 1875/1875 [=====] - 14s 7ms/step - loss: 0.2202 - accuracy: 0.9187
- Epoch 9/10
- 1875/1875 [=====] - 14s 7ms/step - loss: 0.2001 - accuracy: 0.9235
- Epoch 10/10
- 1875/1875 [=====] - 14s 7ms/step - loss: 0.2014 - accuracy: 0.9232
- Testing TF_CNN.
- Classifier algorithm: tf_conv
- Classifier accuracy: 91.940000%
- MNIST D (CNN):
 - Epoch 1/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.2987 - accuracy: 0.9062
 - Epoch 2/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0586 - accuracy: 0.9819
 - Epoch 3/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0456 - accuracy: 0.9866
 - Epoch 4/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0401 - accuracy: 0.9876
 - Epoch 5/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0351 - accuracy: 0.9897
 - Epoch 6/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0285 - accuracy: 0.9910
 - Epoch 7/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0270 - accuracy: 0.9915
 - Epoch 8/10
 - 1875/1875 [=====] - 15s 8ms/step - loss: 0.0251 - accuracy: 0.9917
 - Epoch 9/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0247 - accuracy: 0.9921
 - Epoch 10/10
 - 1875/1875 [=====] - 14s 7ms/step - loss: 0.0235 - accuracy: 0.9927
 - Testing TF_CNN.
 - Classifier algorithm: tf_conv

- Classifier accuracy: 99.320000%
- Cifar 10 output (CNN):
 - Epoch 1/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 1.6969
- accuracy: 0.4087
 - Epoch 2/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 1.1289
- accuracy: 0.5992
 - Epoch 3/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.9683
- accuracy: 0.6585
 - Epoch 4/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.8595
- accuracy: 0.6987
 - Epoch 5/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.7974
- accuracy: 0.7234
 - Epoch 6/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.7396
- accuracy: 0.7382
 - Epoch 7/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.7094
- accuracy: 0.7484
 - Epoch 8/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.6622
- accuracy: 0.7675
 - Epoch 9/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.6363
- accuracy: 0.7776
 - Epoch 10/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 0.6144
- accuracy: 0.7842
 - Testing TF_CNN.
 - Classifier algorithm: tf_conv
 - Classifier accuracy: 72.170000%
- Cifar 100 Fine output (CNN):
 - Epoch 1/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 4.1747
- accuracy: 0.0887
 - Epoch 2/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 3.1656
- accuracy: 0.2321
 - Epoch 3/10

- 1563/1563 [=====] - 16s 10ms/step - loss: 2.7770
- accuracy: 0.2985
- Epoch 4/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.5599
- accuracy: 0.3417
- Epoch 5/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.4305
- accuracy: 0.3657
- Epoch 6/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.3290
- accuracy: 0.3912
- Epoch 7/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.2447
- accuracy: 0.4034
- Epoch 8/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.1698
- accuracy: 0.4246
- Epoch 9/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.1255
- accuracy: 0.4359
- Epoch 10/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 2.0479
- accuracy: 0.4504
- Testing TF_CNN.
- Classifier algorithm: tf_conv
- Classifier accuracy: 45.460000%
- Cifar 100 Coarse output (CNN):
 - Epoch 1/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 2.6582
- accuracy: 0.2115
 - Epoch 2/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 2.0492
- accuracy: 0.3703
 - Epoch 3/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 1.8291
- accuracy: 0.4301
 - Epoch 4/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 1.7122
- accuracy: 0.4608
 - Epoch 5/10
 - 1563/1563 [=====] - 16s 10ms/step - loss: 1.6042
- accuracy: 0.4952
 - Epoch 6/10

- 1563/1563 [=====] - 16s 10ms/step - loss: 1.5359
- accuracy: 0.5145
- Epoch 7/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 1.4711
- accuracy: 0.5327
- Epoch 8/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 1.4239
- accuracy: 0.5495
- Epoch 9/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 1.3842
- accuracy: 0.5625
- Epoch 10/10
- 1563/1563 [=====] - 16s 10ms/step - loss: 1.3377
- accuracy: 0.5745
- Testing TF_CNN.
- Classifier algorithm: tf_conv
- Classifier accuracy: 53.770000%